

Machine Learning

Regularization and Feature Selection

Fabio Vandin

November 25th, 2022

Feature Selection: Scenario

We have a large pool of features

Goal: select a small number of features that will be used by our (final) predictor

Assume $\mathcal{X} = \mathbb{R}^d$.

Goal: learn (final) predictor using $k \ll d$ ^(features) predictors

Feature Selection: Scenario

We have a large pool of features

Goal: select a small number of features that will be used by our (final) predictor

Assume $\mathcal{X} = \mathbb{R}^d$.

Goal: learn (final) predictor using $k \ll d$ predictors

Motivation?

- prevent overfitting: less predictors \Rightarrow hypotheses of lower complexity!
- predictions can be done faster
- useful in many applications!

Feature Selection: Computational Problem

Assume that we use the Empirical Risk Minimization (ERM) procedure.

The problem of selecting k features that minimize the empirical risk can be written as:

Assumption: an hypothesis
best corresponds to
to a vector $\mathbf{w} \in \mathbb{R}^d$
where $\|\mathbf{w}\|_0 = |\{i : w_i \neq 0\}|$

$$\min_{\mathbf{w}} L_S(\mathbf{w}) \text{ subject to } \|\mathbf{w}\|_0 \leq k$$

Feature Selection: Computational Problem

Assume that we use the Empirical Risk Minimization (ERM) procedure.

The problem of selecting k features that minimize the empirical risk can be written as:

$$\min_{\mathbf{w}} L_S(\mathbf{w}) \quad \text{subject to} \quad \|\mathbf{w}\|_0 \leq k$$

where $\|\mathbf{w}\|_0 = |\{i : w_i \neq 0\}|$

How can we solve it?

Subset Selection

How do we find the solution to the problem below?

$$\min_{\mathbf{w}} L_S(\mathbf{w}) \text{ subject to } \|\mathbf{w}\|_0 \leq k$$

Solution 1 :

- i) enumerate all subsets of features with ~~k~~ $\leq k$ elements $\neq \emptyset$
- ii) learn the best model (minimizes the training error) for each of the subsets
- iii) keep the subset of minimum error

Subset Selection

How do we find the solution to the problem below?

$$\min_{\mathbf{w}} L_S(\mathbf{w}) \quad \text{subject to} \quad \|\mathbf{w}\|_0 \leq k$$

Note: the solution will always include k features

Let:

- $\mathcal{I} = \{1, \dots, d\}$;
- given $p = \{i_1, \dots, i_k\} \subseteq \mathcal{I}$: \mathcal{H}_p = hypotheses/models where only features $w_{i_1}, w_{i_2}, \dots, w_{i_k}$ are used

$$P^{(k)} \leftarrow \{J \subseteq \mathcal{I} : |J| = k\};$$

foreach $p \in P^{(k)}$ **do**

$$\quad \left[\begin{array}{l} h_p \leftarrow \arg \min_{h \in \mathcal{H}_p} L_S(h); \end{array} \right.$$

return $h^{(k)} \leftarrow \arg \min_{p \in P^{(k)}} L_S(h_p);$

Complexity? Learn $\Theta\left(\binom{d}{k}\right) \in \Theta(d^k)$ models \Rightarrow exponential algorithm!

\approx it is unlikely that there
is a poly-time algorithm
to solve problem (exactly)

Can we do better?

Proposition

The optimization problem of feature selection NP-hard.

Can we do better?

Proposition

The optimization problem of feature selection NP-hard.

What can we do?

Heuristic solution \Rightarrow greedy algorithms

Greedy Algorithms for Feature Selection

Forward Selection: start from the empty solution, add one feature at the time, until solution has cardinality k

```
 $sol \leftarrow \emptyset;$   
while  $|sol| < k$  do  
  foreach  $i \in \mathcal{I} \setminus sol$  do  
     $p \leftarrow sol \cup \{i\};$   
     $h_p \leftarrow \arg \min_{h \in \mathcal{H}_p} L_S(h);$   
   $sol \leftarrow sol \cup \arg \min_{i \in \mathcal{I} \setminus sol} L_S(h_{sol \cup \{i\}});$   
return  $sol;$ 
```

Complexity? Learns $\Theta(kd)$ models

Backward Selection: start from the solution which includes all features, remove one feature at the time, until solution has cardinality k

Pseudocode: analogous to forward selection [Exercise!]

Complexity? Learns $\Theta((d - k)d)$ models

Notes

We have used only training set to select the best hypothesis...

⇒ we may overfit!

Solution? Use validation! (or cross-validation)

Split data into training data and validation data, learn models on training, evaluate (= pick among different hypothesis models) on validation data. Algorithms are similar.

Note: now the best model (in terms of validation error) may include less than k features!

Subset Selection with Validation Data

S = training data (from data split)

V = validation data (from data split)

Subset Selection with Validation Data

S = training data (from data split)

V = validation data (from data split)

Using training and validation:

```
for  $\ell \leftarrow 0$  to  $k$  do  
   $P^{(\ell)} \leftarrow \{J \subseteq \mathcal{I} : |J| = \ell\};$   
  foreach  $p \in P^{(\ell)}$  do  
     $h_p \leftarrow \arg \min_{h \in \mathcal{H}_p} L_S(h);$   
   $h^{(\ell)} \leftarrow \arg \min_{p \in P^{(\ell)}} L_V(h_p);$   
return  $\arg \min_{h \in \{h^{(0)}, h^{(1)}, \dots, h^{(k)}\}} L_V(h)$ 
```

Forward Selection with Validation Data

Using training and validation:

```
sol  $\leftarrow \emptyset$ ;  
while  $|sol| < k$  do  
  foreach  $i \in \mathcal{I} \setminus sol$  do  
     $p \leftarrow sol \cup \{i\}$ ;  
     $h_p \leftarrow \arg \min_{h \in \mathcal{H}_p} L_S(h)$ ;  
   $sol \leftarrow sol \cup \arg \min_{i \in \mathcal{I} \setminus sol} L_V(h_{sol \cup \{i\}})$ ;  
return sol;
```

Backward Selection with validation: similar [Exercise]

Similar approach for all algorithms with cross-validation [Exercise]

Bibliography [UML]

Regularization and Ridge Regression: Chapter 12

- no Section 13.3;
- Section 13.4 only up to Corollary 13.8 (excluded)

Feature Selection and LASSO: Chapter 25

- only Section 25.1.2 (introduction and “Backward Elimination”) and 25.1.3