

Machine Learning

Neural Networks

Fabio Vandin


December 6th, 2022



Neuron

Neuron: function $\mathbf{x} \rightarrow \sigma(\langle \mathbf{v}, \mathbf{x} \rangle)$, with $\mathbf{x} \in \mathbb{R}^d$

$\sigma : \mathbb{R} \rightarrow \mathbb{R}$ is the *activation function*


Example: \mathbb{R}^5





We will consider σ to be one among:


- sign function: $\sigma(a) = \text{sign}(a)$
- threshold function: $\sigma(a) = \mathbb{1}[a > 0]$
- sigmoid function: $\sigma(a) = \frac{1}{1+e^{-a}}$



Neural Network (NN)

Obtained by connecting many neurons together.

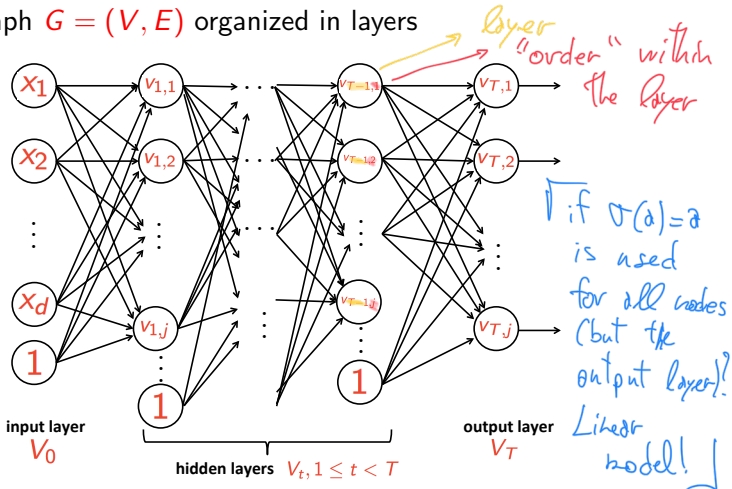
We focus on **feedforward neural networks**, defined by a directed acyclic graph $G = (V, E)$ organized in layers



Neural Network (NN)

Obtained by connecting many neurons together.

We focus on **feedforward neural networks**, defined by a directed acyclic graph $G = (V, E)$ organized in layers




Each edge e has a **weight** $w(e)$ specified by $w : E \rightarrow \mathbb{R}$

Point of View of One Node

Consider node $v_{t+1,j}$, $0 \leq t < T$. Let


layer $t+1$, j -th node in the layer



Point of View of One Node

Consider node $v_{t+1,j}$, $0 \leq t < T$. Let


- $a_{t+1,j}(x)$: its *input* when x is fed to the NN before $\nabla()$ is applied



Point of View of One Node

Consider node $v_{t+1,j}$, $0 \leq t < T$. Let


- $a_{t+1,j}(\mathbf{x})$: its *input* when \mathbf{x} is fed to the NN
- $o_{t+1,j}(\mathbf{x})$: its *output* when \mathbf{x} is fed to the NN



Point of View of One Node

Consider node $v_{t+1,j}$, $0 \leq t < T$. Let

- $a_{t+1,j}(\mathbf{x})$: its *input* when \mathbf{x} is fed to the NN
- $o_{t+1,j}(\mathbf{x})$: its *output* when \mathbf{x} is fed to the NN




Then:
$$a_{t+1,j}(\mathbf{x}) = \sum_{r:(v_{t,r}, v_{t+1,j}) \in E} w((v_{t,r}, v_{t+1,j})) o_{t,r}(\mathbf{x})$$

Point of View of One Node

Consider node $v_{t+1,j}$, $0 \leq t < T$. Let

- $a_{t+1,j}(\mathbf{x})$: its *input* when \mathbf{x} is fed to the NN
- $o_{t+1,j}(\mathbf{x})$: its *output* when \mathbf{x} is fed to the NN



If $\nabla(o) = 0$
 \Rightarrow neuron
corresponds
to a linear
model


$$\text{Then: } a_{t+1,j}(\mathbf{x}) = \sum_{r: (v_{t,r}, v_{t+1,j}) \in E} w((v_{t,r}, v_{t+1,j})) o_{t,r}(\mathbf{x})$$

$$o_{t+1,j}(\mathbf{x}) = \sigma(a_{t+1,j}(\mathbf{x}))$$

Neural Network: Formalism

Neural network: described by directed acyclic graph $G = (V, E)$ and weight function $w : E \rightarrow \mathbb{R}$

- $V = \bigcup_{t=0}^T V_t, V_i \cap V_j = \emptyset \quad \forall i \neq j$
- $e \in E$ can only go from V_t to V_{t+1} for some t
- $V_0 = \text{input layer}$
- $V_T = \text{output layer}$
- $V_t, 0 < t < T = \text{hidden layers}$
- $T = \text{depth}$
- $|V| = \text{size of the network}$
- $\max_t |V_t| = \text{width of the network}$



depth?

3 ? 4.25
2 ? 33.25
1 ? 12.25
0 ? 0.25

Neural Network: Formalism

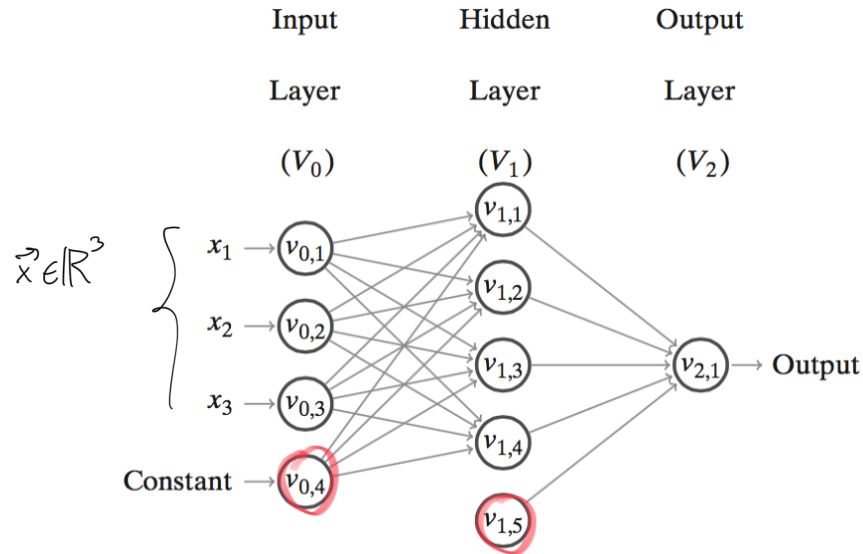
Neural network: described by directed acyclic graph $G = (V, E)$ and weight function $w : E \rightarrow \mathbb{R}$

- $V = \cup_{t=0}^T V_t, V_i \cap V_j = \emptyset \ \forall i \neq j$
- $e \in E$ can only go from V_t to V_{t+1} for some t
- $V_0 = \text{input layer}$
- $V_T = \text{output layer}$
- $V_t, 0 < t < T = \text{hidden layers}$
- $T = \text{depth}$
- $|V| = \text{size of the network}$
- $\max_t |V_t| = \text{width of the network}$

Notes:

- for binary classification and regression (1 variable): output layer has 1 node
- different layers could have different activation functions (e.g., output layer)

Example



depth = 2, size = 10, width = 5

Exercise

Assume that for each node the activation function $\sigma(z) : \mathbb{R} \rightarrow \mathbb{R}$ is defined as

$$\sigma(z) = \begin{cases} 1 & z \geq 1 \\ z & -1 \leq z < 1 \\ -1 & z < -1 \end{cases}$$

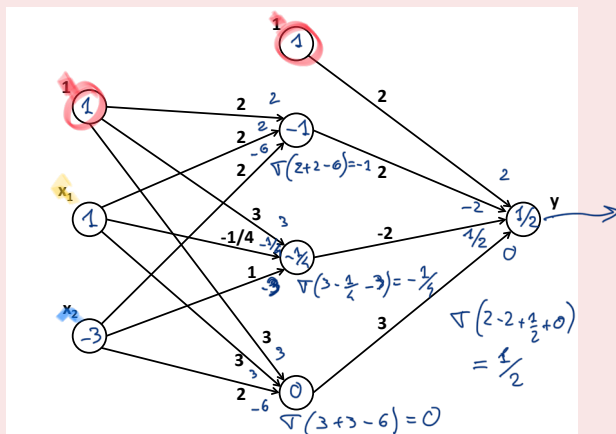
and consider the neural network in the next slide, compute the value of the output y when the input $\mathbf{x} \in \mathbb{R}^2$ is

$$\mathbf{x} = [1 \quad -3]^\top$$

Exercise (continue)

$$\vec{x} = \begin{bmatrix} 1 \\ -3 \end{bmatrix}$$

$$= \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$



\Rightarrow predicted value for $\vec{x} = \begin{bmatrix} 1 \\ -3 \end{bmatrix}$ is $\frac{1}{2}$.

Hypothesis Set of a NN

Architecture of a NN: (V, E, σ)

Once we specify the architecture and w , we obtain a function:

$$h_{V,E,\sigma,w} : \mathbb{R}^{|V_0|-1} \rightarrow \mathbb{R}^{|V_T|}$$

architecture \rightarrow *weights (on edges)*

Hypothesis Set of a NN

Architecture of a NN: (V, E, σ)

Once we specify the architecture and w , we obtain a function:

$$h_{V,E,\sigma,w} : \mathbb{R}^{|V_0|-1} \rightarrow \mathbb{R}^{|V_T|}$$

The *hypothesis class* of a neural network is defined by *fixing* its architecture:

$$\mathcal{H}_{\underbrace{V,E,\sigma}_{\text{architecture}}} = \{h_{V,E,\sigma,w} : w \text{ is a mapping from } E \text{ to } \mathbb{R}\}$$

Hypothesis Set of a NN

Architecture of a NN: (V, E, σ)

Once we specify the architecture and w , we obtain a function:

$$h_{V,E,\sigma,w} : \mathbb{R}^{|V_0|-1} \rightarrow \mathbb{R}^{|V_T|}$$

The *hypothesis class* of a neural network is defined by *fixing* its architecture:


$$\mathcal{H}_{V,E,\sigma} = \{h_{V,E,\sigma,w} : w \text{ is a mapping from } E \text{ to } \mathbb{R}\}$$

Question: what type of functions can be implemented using a neural network?

Exercise (expressiveness of NNs)


Let $\mathbf{x} = [x_1, x_2] \in \{-1, 1\}^2$, and let the training data be represented by the following table:

x_1	x_2	y
-1	-1	-1
-1	1	1
1	-1	1
1	1	-1




Consider the NN in the figure above, where the activation function for each hidden node and the output node is the *sign* function. Assume that the network's weights are constrained to be in $\{-1, 1\}$.

- 1 Find network's weights so that the training error is 0.
- 2 Use example above to motivate the fact that NNs are *richer* models than linear models.



2) The training set represents the "XOR" function, that is not linear



$$\vec{x} \in \{-1, 1\}^2$$

Such dataset cannot be perfectly classified with a linear model, but NNs can perfectly classify the data \Rightarrow NNs are richer models than linear models!

General construction: let's take an arbitrary function
 $f: \{-1, 1\}^d \rightarrow \{-1, 1\}$.

Goal: build a NN that corresponds to f : if the input is \vec{x} , then the prediction of such NN is $f(\vec{x})$

i) consider \vec{x} such that $f(\vec{x}) = 1$: for each such \vec{x} , there is a neuron in the (only) hidden layer that "corresponds" to \vec{x} . The neuron implements:

$$g_i(\vec{x}') = \text{sign}(\langle \vec{x}, \vec{x}' \rangle - d + 1)$$


weights on the
incoming edges

input to the
neuron


neuron $g_i(\vec{x}^1)$

"corresponds" to


$$\vec{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$



- output node: "implements" $h(\vec{x}^1) = \text{sign}\left(\sum_{i=1}^k g_i(\vec{x}^1) + k-1\right)$
 where $k = \#$ of inputs \vec{x} s.t. $f(\vec{x}) = 1$



The overall net is :



Exercise: show that the net "computes" $f()$.

Expressiveness of NN

set of hypotheses for the \mathcal{H} where the architecture is given by (V, E) and the sign activation function

Proposition

For every d , there exists a graph (V, E) of depth 2 such that $\mathcal{H}_{V, E, \text{sign}}$ contains all functions from $\{-1, 1\}^d$ to $\{-1, 1\}$

Expressiveness of NN

Proposition

For every d , there exists a graph (V, E) of depth 2 such that $\mathcal{H}_{V,E,\text{sign}}$ contains all functions from $\{-1, 1\}^d$ to $\{-1, 1\}$

NN can implement every boolean function!

Unfortunately the graph (V, E) is very big...

Proposition

For every d , let $s(d)$ be the minimal integer such that there exists a graph (V, E) with $|V| = s(d)$ such that $\mathcal{H}_{V,E,\text{sign}}$ contains all functions from $\{-1, 1\}^d$ to $\{-1, 1\}$. Then $s(d)$ is an exponential function of d .

Note: similar result for $\sigma = \text{sigmoid}$

Proposition

For every fixed $\varepsilon > 0$ and every Lipschitz function $f : [-1, 1]^d \rightarrow [-1, 1]$ it is possible to construct a neural network such that for every input $\mathbf{x} \in [-1, 1]^d$ the output of the neural network is in $[f(\mathbf{x}) - \varepsilon, f(\mathbf{x}) + \varepsilon]$.

Note: first result proved by Cybenko (1989) for sigmoid activation function, requires only 1 hidden layer!

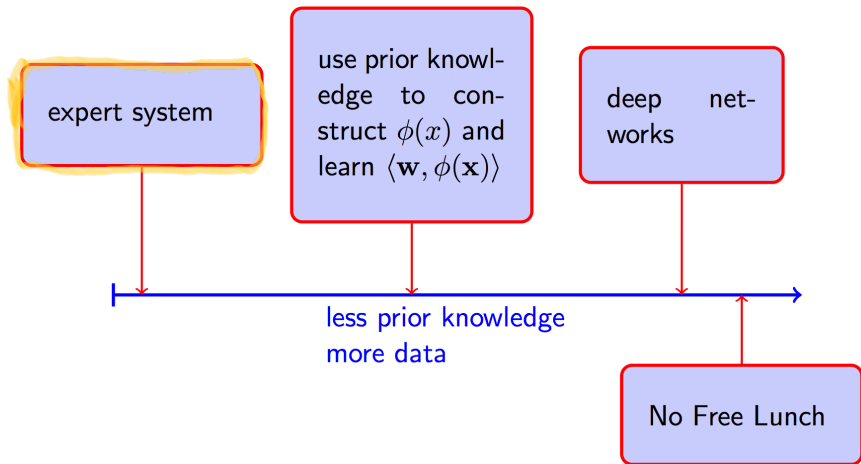
NNs are **universal approximators!**

But again...

Proposition

Fix some $\varepsilon \in (0, 1)$. For every d , let $s(d)$ be the minimal integer such that there exists a graph (V, E) with $|V| = s(d)$ such that $\mathcal{H}_{V,E,\sigma}$, with $\sigma = \text{sigmoid}$, can approximate, with precision ε , every 1-Lipschitz function $f : [-1, 1]^d \rightarrow [-1, 1]$. Then $s(d)$ is exponential in d .

An Extremely Powerful Hypothesis Class...



Sample Complexity of NNs

How much data is needed to learn with NNs?

Proposition

The VC dimension of $\mathcal{H}_{V,E,\text{sign}} = O(|E| \log |E|)$

Different σ ?

Proposition

Let σ be the sigmoid function. The VC dimension of $\mathcal{H}_{V,E,\sigma}$ is:

- $\Omega(|E|^2)$
- $O(|V|^2|E|^2)$

\Rightarrow large NNs require a lot of data!

Question: assume we have a lot of data, can we find the best hypothesis?

Runtime of Learning NNs

Informally: applying the ERM rule with respect to $\mathcal{H}_{V,E,\text{sign}}$ is *computationally difficult*, even for small NN...

Proposition

Let $k \geq 3$. For every d , let (V, E) be a layered graph with d input nodes, $k + 1$ nodes at the (only) hidden layer, where one of them is the constant neuron, and a single output node. Then, it is NP-hard to implement the ERM rule with respect to $\mathcal{H}_{V,E,\text{sign}}$.

Well maybe the above is only for very specific cases...

- instead of ERM rule, find h close to ERM? **Computationally infeasible!** (probably)
- other activation functions (e.g., sigmoid)? **Computationally infeasible!** (probably)
- smart embedding in larger network? **Computationally infeasible!** (probably)